

GitHub: https://github.com/Muralikrishna9550/Assignment_4.git

Video: https://drive.google.com/file/d/1a1N_NNaR2uTBuCHiiBIBD-x42EorfMeu/view?usp=drive_link

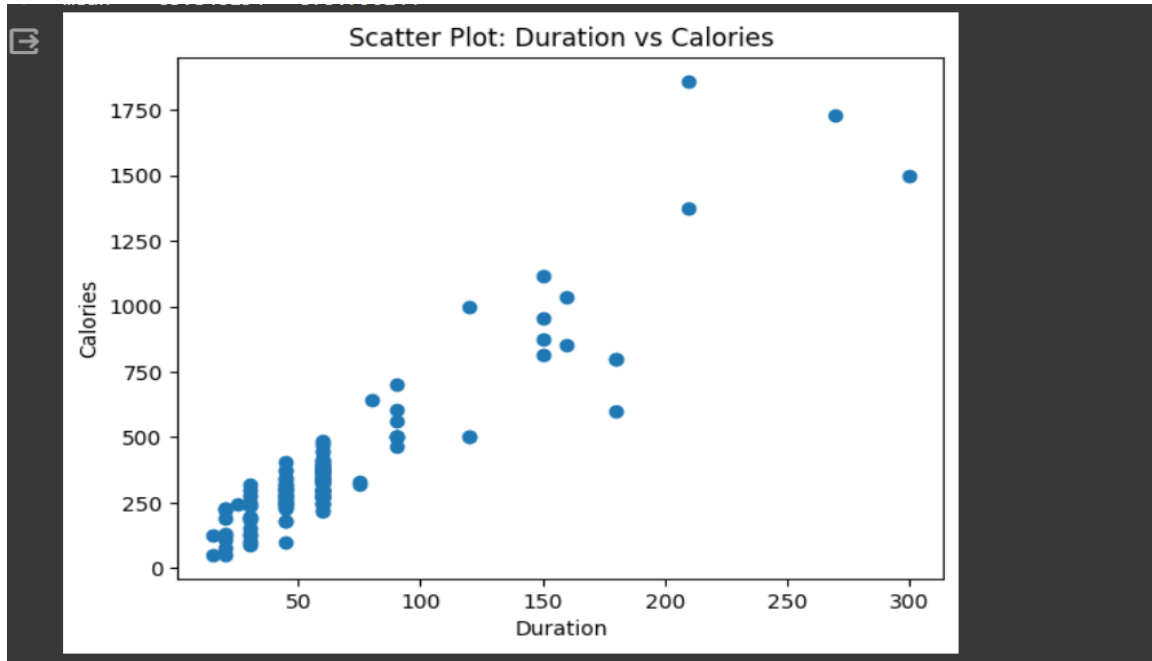
1)

```
import pandas as pd
data = pd.read_csv('data.csv')
description = data.describe()
print(description)
null_values = data.isnull().sum()
print("Null Values:")
print(null_values)
data.fillna(data.mean(), inplace=True)
agg_columns = ['Duration', 'Calories']
agg_functions = {
    'Duration': ['min', 'max', 'count', 'mean'],
    'Calories': ['min', 'max', 'count', 'mean']
}
aggregated_data = data[agg_columns].agg(agg_functions)
print("Aggregated Data:")
print(aggregated_data)
filtered_data1 = data[(data['Calories'] >= 500) & (data['Calories'] <= 1000)]
filtered_data2 = data[(data['Calories'] > 500) & (data['Pulse'] < 100)]
df_modified = data.drop(columns=['Maxpulse'])
data.drop(columns=['Maxpulse'], inplace=True)
data['Calories'] = data['Calories'].astype(int)
import matplotlib.pyplot as plt
plt.scatter(data['Duration'], data['Calories'])
plt.xlabel('Duration')
plt.ylabel('Calories')
plt.title('Scatter Plot: Duration vs Calories')
plt.show()
```

OutPut:

→

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000
Null Values:				
Duration	0			
Pulse	0			
Maxpulse	0			
Calories	5			
dtype: int64				
Aggregated Data:				
	Duration	Calories		
min	15.000000	50.300000		
max	300.000000	1860.400000		
count	169.000000	169.000000		
mean	63.846154	375.790244		



2)

```
[ ] import pandas as pd
    file_name = 'Salary_Data.csv'
    data = pd.read_csv(file_name)
    print(data.head())
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
file_name = 'Salary_Data.csv'
data = pd.read_csv(file_name)
X = data[['YearsExperience']]
y = data['Salary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/3, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
plt.scatter(X_train, y_train, color='blue', label='Training Data')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Training Data - Salary vs. Years of Experience')
plt.legend()
plt.show()
plt.scatter(X_test, y_test, color='red', label='Test Data')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Test Data - Salary vs. Years of Experience')
plt.legend()
plt.show()
```

Output:

